

Class 11: Transcriptomics and the analysis of RNA-Seq data

Taylor F. (A59010460)

2/23/2022

Installing DESeq2

```
#Do this in the console in the future
install.packages("BiocManager")
BiocManager::install()

#We will also need the DESeq2 package
BiocManager::install("DESeq2")

#NOTE: Answer NO to prompts to install from source or update

#Run library(DESeq2) in the console
```

Today we will run differential expression analysis of some published data from Himes et al. where the authors used a published RNA-seq experiment where airway smooth muscle cells were treated with dexamethasone, a synthetic glucocorticoid steroid with anti-inflammatory effects.

Importing countData and colData

```
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
metadata <- read.csv("airway_metadata.csv")

#Preview the counts dataset
head(counts)
```

```
##                SRR1039508 SRR1039509 SRR1039512 SRR1039513 SRR1039516
## ENSG000000000003         723         486         904         445         1170
## ENSG000000000005          0          0          0          0          0
## ENSG000000000419         467         523         616         371         582
## ENSG000000000457         347         258         364         237         318
## ENSG000000000460          96          81          73          66         118
## ENSG000000000938          0          0          1          0          2
##                SRR1039517 SRR1039520 SRR1039521
## ENSG000000000003        1097         806         604
## ENSG000000000005          0          0          0
## ENSG000000000419         781         417         509
```

```
## ENSG00000000457      447      330      324
## ENSG00000000460       94      102       74
## ENSG00000000938        0        0        0
```

```
#Determine how many genes are in this counts dataset
nrow(counts)
```

```
## [1] 38694
```

```
#Determine how many control cell lines used
ncol(counts)
```

```
## [1] 8
```

```
#Look at the metadata dataset
metadata
```

```
##           id      dex celltype      geo_id
## 1 SRR1039508 control   N61311 GSM1275862
## 2 SRR1039509 treated   N61311 GSM1275863
## 3 SRR1039512 control   N052611 GSM1275866
## 4 SRR1039513 treated   N052611 GSM1275867
## 5 SRR1039516 control   N080611 GSM1275870
## 6 SRR1039517 treated   N080611 GSM1275871
## 7 SRR1039520 control   N061011 GSM1275874
## 8 SRR1039521 treated   N061011 GSM1275875
```

Question 1:

There are 38694 rows, which translate to genes, in the counts dataset.

Question 2:

There are 8 rows, which translate to individual cell lines, in the counts dataset.

Based on the metadata, it looks like we have four drug-treated and four control cell lines. Our first question is does the drug do anything?

First, we want to check if the metadata matches the counts data order:

```
#Grab the id column values of the metadata dataset
metadata$id
```

```
## [1] "SRR1039508" "SRR1039509" "SRR1039512" "SRR1039513" "SRR1039516"
## [6] "SRR1039517" "SRR1039520" "SRR1039521"
```

```
#Grab the column names of the counts dataset
colnames(counts)
```

```
## [1] "SRR1039508" "SRR1039509" "SRR1039512" "SRR1039513" "SRR1039516"
## [6] "SRR1039517" "SRR1039520" "SRR1039521"
```

```
#Ask if these values are equivalent (all returned values should be TRUE)
metadata$id == colnames(counts)
```

```
## [1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```

```
#Alternative method
all(metadata$id == colnames(counts))
```

```
## [1] TRUE
```

```
#Fancy method
if(all(metadata$id == colnames(counts))) {cat("Let's do this!")}
```

```
## Let's do this!
```

Next, we want to separate the two conditions (control and treated) from “counts” and use a summary statistic to make comparison easier.

```
#Control first
#We need to find the ID associated with control conditions (i.e. which id columns also have dex == "control")
control.inds <- metadata[metadata$dex == "control", "id"]
control.inds
```

```
## [1] "SRR1039508" "SRR1039512" "SRR1039516" "SRR1039520"
```

```
#We now use these indices to search through counts and finds these columns
control.counts <- counts[,control.inds]
head(control.counts)
```

```
##           SRR1039508 SRR1039512 SRR1039516 SRR1039520
## ENSG00000000003      723       904       1170       806
## ENSG00000000005         0         0         0         0
## ENSG00000000419      467       616       582       417
## ENSG00000000457      347       364       318       330
## ENSG00000000460       96        73       118       102
## ENSG00000000938         0         1         2         0
```

```
#Question 4: We can now access the treated values, too
treated.inds <- metadata[metadata$dex == "treated", "id"]
treated.inds
```

```
## [1] "SRR1039509" "SRR1039513" "SRR1039517" "SRR1039521"
```

```
treated.counts <- counts[,treated.inds]
head(treated.counts)
```

```
##           SRR1039509 SRR1039513 SRR1039517 SRR1039521
## ENSG00000000003      486       445       1097       604
## ENSG00000000005         0         0         0         0
## ENSG00000000419      523       371       781       509
## ENSG00000000457      258       237       447       324
## ENSG00000000460       81        66        94        74
## ENSG00000000938         0         0         0         0
```

Find the mean count value for each row (i.e. gene). We could use the 'apply()' function or more simply the 'rowMeans()' function.

```
#Let's now find the means for each of these groups
```

```
control.mean <- rowMeans(control.counts)
```

```
treated.mean <- rowMeans(treated.counts)
```

```
#Question 5a: We can plot these mean values against one another
```


```
plot(treated.mean, control.mean, log = "xy")
```

```
## Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 x values <= 0 omitted  
## from logarithmic plot
```

```
## Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 y values <= 0 omitted  
## from logarithmic plot
```

```
#Question 5b: We can also use ggplot
```


```
library(ggplot2)
```



```
ggplot(counts, aes(treated.mean, control.mean)) + geom_point() + scale_y_continuous(trans='log10') + scale_x_continuous(trans='log10')
```

```
## Warning: Transformation introduced infinite values in continuous y-axis
```

```
## Warning: Transformation introduced infinite values in continuous x-axis
```



#Question 6: Try plotting both axes on a log scale. What is the argument to plot() that allows you to do this?
#log = "xy"

We often use log 2 transformation because it is easier to understand, visually.

```
#20/20
#log2(20/20)
#log2(40/20)
#log2(10/20)
#log2(80/20)

#Store the log2 fold change between treated and control groups
log2fc <- log2(treated.mean/control.mean)
```

Finding and filtering zero values

```
#We need to find and remove the genes that have zeros for values
log2fc[1:6]
```

```
## ENSG00000000003 ENSG00000000005 ENSG000000000419 ENSG000000000457 ENSG000000000460
##      -0.45303916          NaN      0.06900279      -0.10226805      -0.30441833
## ENSG000000000938
##      -Inf
```

```
meancounts <- data.frame(control.mean, treated.mean, log2fc)
head(meancounts[,1:2] == 0)
```

```
##           control.mean treated.mean
## ENSG000000000003      FALSE      FALSE
## ENSG000000000005       TRUE       TRUE
## ENSG000000000419      FALSE      FALSE
## ENSG000000000457      FALSE      FALSE
## ENSG000000000460      FALSE      FALSE
## ENSG000000000938      FALSE       TRUE
```

```
z <- data.frame(x = c(1,2,0,4), y = c(1,2,0,0))
```

```
#Report which indices are TRUE and FALSE (i.e. sum to 0 or greater than 0) and gives index information
which(z == 0, arr.ind = TRUE)
```

```
##      row col
## [1,]   3   1
## [2,]   3   2
## [3,]   4   2
```

```
unique(which(z == 0, arr.ind = TRUE)[,"row"])
```

```
## [1] 3 4
```

```
#Apply this principle to the meancounts dataset
to.rm <- sort(unique(which(meancounts[,1:2] == 0, arr.ind = TRUE)[,"row"])))
mycounts <- meancounts[-to.rm,]
#mycounts
```

There are 21817 genes left over after removing zero values.

How many genes have a log2fc more than +2 (i.e. upregulated)?

```
up.ind <- mycounts$log2fc > 2
down.ind <- mycounts$log2fc < (-2)
sum(up.ind)
```

```
## [1] 250
```

```
sum(down.ind)
```

```
## [1] 367
```

There are 250 genes upregulated by at least 2log() and rsum(down.ind) genes downregulated by at least 2log().

But this approach does not treat fold-changes equally. There may be significant fold change increases or decreases in a gene expressed very low, and a comparatively large change that is not included in another gene that has a higher basal level of expression.

DESeq2

This approach will be the right way and will give us the stats.

```
library(DESeq2)
```

```
## Loading required package: S4Vectors
```

```
## Loading required package: stats4
```

```
## Loading required package: BiocGenerics
```

```
##
```

```
## Attaching package: 'BiocGenerics'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      IQR, mad, sd, var, xtabs
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      anyDuplicated, append, as.data.frame, basename, cbind, colnames,  
##      dirname, do.call, duplicated, eval, evalq, Filter, Find, get, grep,  
##      grepl, intersect, is.unsorted, lapply, Map, mapply, match, mget,  
##      order, paste, pmax, pmax.int, pmin, pmin.int, Position, rank,  
##      rbind, Reduce, rownames, sapply, setdiff, sort, table, tapply,  
##      union, unique, unsplit, which.max, which.min
```

```
##
```

```
## Attaching package: 'S4Vectors'
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      expand.grid, I, unname
```

```
## Loading required package: IRanges
```

```
## Loading required package: GenomicRanges
```

```
## Loading required package: GenomeInfoDb
```

```
## Loading required package: SummarizedExperiment
```

```
## Loading required package: MatrixGenerics
```

```
## Loading required package: matrixStats
```

```
##
```

```
## Attaching package: 'MatrixGenerics'
```

```
## The following objects are masked from 'package:matrixStats':
```

```
##
```

```
## colAlls, colAnyNAs, colAnys, colAvgPerRowSet, colCollapse,  
## colCounts, colCummaxs, colCummins, colCumprods, colCumsums,  
## colDiffs, colIQRDiffs, colIQRs, colLogSumExps, colMadDiffs,  
## colMads, colMaxs, colMeans2, colMedians, colMins, colOrderStats,  
## colProds, colQuantiles, colRanges, colRanks, colSdDiffs, colSds,  
## colSums2, colTabulates, colVarDiffs, colVars, colWeightedMads,  
## colWeightedMeans, colWeightedMedians, colWeightedSds,  
## colWeightedVars, rowAlls, rowAnyNAs, rowAnys, rowAvgPerColSet,  
## rowCollapse, rowCounts, rowCummaxs, rowCummins, rowCumprods,  
## rowCumsums, rowDiffs, rowIQRDiffs, rowIQRs, rowLogSumExps,  
## rowMadDiffs, rowMads, rowMaxs, rowMeans2, rowMedians, rowMins,  
## rowOrderStats, rowProds, rowQuantiles, rowRanges, rowRanks,  
## rowSdDiffs, rowSds, rowSums2, rowTabulates, rowVarDiffs, rowVars,  
## rowWeightedMads, rowWeightedMeans, rowWeightedMedians,  
## rowWeightedSds, rowWeightedVars
```

```
## Loading required package: Biobase
```

```
## Welcome to Bioconductor
```

```
##
```

```
## Vignettes contain introductory material; view with  
## 'browseVignettes()'. To cite Bioconductor, see  
## 'citation("Biobase")', and for packages 'citation("pkgname")'.
```

```
##
```

```
## Attaching package: 'Biobase'
```

```
## The following object is masked from 'package:MatrixGenerics':
```

```
##
```

```
## rowMedians
```

```
## The following objects are masked from 'package:matrixStats':
```

```
##
```

```
## anyMissing, rowMedians
```

```
#Let's look at metadata again and set up the object that DESeq needs with the 'DESeqDataSetFromMatrix()'  
dds <- DESeqDataSetFromMatrix(countData = counts, colData = metadata, design =~ dex)
```

```
## converting counts to integer mode
```

```
## Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in  
## design formula are characters, converting to factors
```



```
dds
```

```
## class: DESeqDataSet
## dim: 38694 8
## metadata(1): version
## assays(1): counts
## rownames(38694): ENSG000000000003 ENSG000000000005 ... ENSG00000283120
## ENSG00000283123
## rowData names(0):
## colnames(8): SRR1039508 SRR1039509 ... SRR1039520 SRR1039521
## colData names(4): id dex celltype geo_id
```

```
dds <- DESeq(dds)
```

```
## estimating size factors
```

```
## estimating dispersions
```

```
## gene-wise dispersion estimates
```

```
## mean-dispersion relationship
```

```
## final dispersion estimates
```

```
## fitting model and testing
```

```
res <- results(dds)
res
```

```
## log2 fold change (MLE): dex treated vs control
## Wald test p-value: dex treated vs control
## DataFrame with 38694 rows and 6 columns
##      baseMean log2FoldChange      lfcSE      stat      pvalue
##      <numeric>      <numeric> <numeric> <numeric> <numeric>
## ENSG000000000003  747.1942    -0.3507030  0.168246 -2.084470  0.0371175
## ENSG000000000005    0.0000         NA         NA         NA         NA
## ENSG000000000419  520.1342     0.2061078  0.101059  2.039475  0.0414026
## ENSG000000000457  322.6648     0.0245269  0.145145  0.168982  0.8658106
## ENSG000000000460   87.6826    -0.1471420  0.257007 -0.572521  0.5669691
## ...           ...           ...           ...           ...
## ENSG00000283115   0.000000         NA         NA         NA         NA
## ENSG00000283116   0.000000         NA         NA         NA         NA
## ENSG00000283119   0.000000         NA         NA         NA         NA
## ENSG00000283120   0.974916    -0.668258    1.69456 -0.394354  0.693319
## ENSG00000283123   0.000000         NA         NA         NA         NA
##      padj
##      <numeric>
## ENSG000000000003  0.163035
## ENSG000000000005    NA
## ENSG000000000419  0.176032
## ENSG000000000457  0.961694
```

```
## ENSG00000000460 0.815849
## ...
## ENSG00000283115 NA
## ENSG00000283116 NA
## ENSG00000283119 NA
## ENSG00000283120 NA
## ENSG00000283123 NA
```

A main result figure

A common main result figure from this type of analysis is called a volcano plot. This is a plot of log2 fold change on the x axis vs. p-value.

```
#We can also color by significance and fold change
sigcols <- rep("gray", nrow(res))
sigcols[ abs(res$log2FoldChange) > 2 & (res$padj) <= 0.05 ] <- "darkgreen"
sigcols[ abs(res$log2FoldChange) < 2 | (res$padj) > 0.05 ] <- "red"

#Plot this data
plot(res$log2FoldChange, -log(res$padj), xlab = "Log2(FoldChange)",
      ylab = "-Log(P-value)", col = sigcols)

#Add significance lines to the plot
abline(v = c(-2,2), col = "black", lty = 2)
abline(h = -log(0.05), col = "black", lty = 2)
```

